

Tasking with OpenMP

- 1 Programming Parallel Shared Memory Computers
 - tasking with OpenMP
- 2 Parallel Recursive Functions
 - the Fibonacci numbers
 - parallel recursive quadrature
- 3 Bernstein's Conditions
 - dependency analysis
- 4 Task Dependences
 - flow dependence
 - parallel blocked matrix multiplication

MCS 572 Lecture 15
Introduction to Supercomputing
Jan Verschelde, 15 February 2021

Tasking with OpenMP

1 Programming Parallel Shared Memory Computers

- tasking with OpenMP

2 Parallel Recursive Functions

- the Fibonacci numbers
- parallel recursive quadrature

3 Bernstein's Conditions

- dependency analysis

4 Task Dependences

- flow dependence
- parallel blocked matrix multiplication

tasking with OpenMP

The OpenMP Application Programming Interface provides a model for parallel programming that is portable across shared memory architectures from different vendors.

The `gcc` compiler supports the OpenMP API, via `gcc -fopenmp`.

Two reference documents for this lecture:

- the OpenMP API Specification, and
- the OpenMP API Examples.

Tasking constructs provide units of work to a thread for execution.

Tasking with OpenMP

1 Programming Parallel Shared Memory Computers

- tasking with OpenMP

2 Parallel Recursive Functions

- the Fibonacci numbers
- parallel recursive quadrature

3 Bernstein's Conditions

- dependency analysis

4 Task Dependences

- flow dependence
- parallel blocked matrix multiplication

the Fibonacci numbers

The sequence of Fibonacci numbers F_n are defined as

$$F_0 = 0, \quad F_1 = 1, \quad \text{and for } n > 1 : F_n = F_{n-1} + F_{n-2}.$$

This leads to a natural recursive function.

- 1 The recursion generates many function calls.
- 2 While inefficient to compute F_n ,
this recursion serves as a parallel pattern.

The parallel version is part of the
OpenMP Application Programming Interface Examples.

a parallel recursive Fibonacci function

The Fibonacci function with tasking

- demonstrates the generation of a large number of tasks with one thread.
- No parallelism will result from this example.

But it is instructive to introduce basic task constructs.

- The `task` construct defines an explicit task.
- The `taskwait` construct synchronizes sibling tasks.
- The `shared` clause of a task construct declares a variable to be shared by tasks.

command line arguments and number of threads

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int fib ( int n );
/* Returns the n-th Fibonacci number,
 * computed recursively with tasking. */

int main ( int argc, char *argv[] )
{
    int n;

    if(argc > 1)
        n = atoi(argv[1]);
    else
    {
        printf("Give n : "); scanf("%d", &n);
    }
    omp_set_num_threads(8);
```

the parallel region and single construct

The program from the previous slide continues:

```
#pragma omp parallel
{
    #pragma omp single
    printf("F(%d) = %d\n", n, fib(n));
}
```

The `single` construct specifies that the statement is executed by only one thread in the team.

In this example, one thread generates many tasks.

a parallel recursive Fibonacci function

```
int fib ( int n )
{
    if(n < 2)
        return n;
    else
    {
        int left,right;    // shared by all tasks

        #pragma omp task shared(left)
        left = fib(n-1);

        #pragma omp task shared(right)
        right = fib(n-2);

                                // synchronize tasks
        #pragma omp taskwait
        return left + right;
    }
}
```

Tasking with OpenMP

1 Programming Parallel Shared Memory Computers

- tasking with OpenMP

2 Parallel Recursive Functions

- the Fibonacci numbers
- parallel recursive quadrature

3 Bernstein's Conditions

- dependency analysis

4 Task Dependences

- flow dependence
- parallel blocked matrix multiplication

parallel recursive quadrature

Apply a numerical integration rule $R(f, a, b, n)$ to $\int_a^b f(x)dx$.

The rule $R(f, a, b, n)$ takes on input

- the function f , bounds a, b of $[a, b]$, and
- the number n of function evaluations.

The rule returns an approximation A and an error estimate e .

If e is larger than some tolerance, then

- 1 $c = (b - a)/2$ is the middle of $[a, b]$,
- 2 compute $A_1, e_1 = R(f, a, c, n)$,
- 3 compute $A_2, e_2 = R(f, c, b, n)$,
- 4 return $A_1 + A_2, e_1 + e_2$.

This is the same pattern as Fibonacci.

the composite Trapezoidal rule applied recursively

Using n subintervals of $[a, b]$, the rule is

$$R(f, a, b, n) = \frac{h}{2}(f(a) + f(b)) + h \sum_{i=1}^{n-1} f(a + ih), \quad h = \frac{b - a}{n}.$$

Our setup: $f(x) = e^x$, $[a, b] = [0, 1]$, $\int_0^1 e^x dx = e - 1$.

Keep n fixed. Let d be the depth of the recursion. The level is ℓ .

$\mathcal{F}(\ell, d, f, a, b, n)$:

 If $\ell = d$ then

 return $R(f, a, b, n)$


 else

$c = (b - a)/2$

 return $\mathcal{F}(\ell+1, d, f, a, c, n) + \mathcal{F}(\ell+1, d, f, c, b, n)$.

the tree of function calls

The root of the tree is the first call, omitting the value for n .



At the leaves, the rule is applied.

As all computations are concentrated at the leaves,
we expect speedups from a parallel execution.

a recursive parallel integration function

```
double rectraprule
( int level, int depth,
  double (*f) ( double x ), double a, double b, int n )
{
    if(level == depth)
        return traprule(f,a,b,n);
    else
    {
        double middle = (b - a)/2;
        double left,right;

        #pragma omp task shared(left)
        left = rectraprule(level+1,depth,f,a,middle,n);

        #pragma omp task shared(right)
        right = rectraprule(level+1,depth,f,middle,b,n);

        #pragma omp taskwait
        return left + right;
    }
}
```

with 8 threads, with OpenMP 4.5 on pascal

```
$ time ./comptraprec 200000 10
approximation = 1.7182818284620265e+00
exp(1) - 1 = 1.7182818284590451e+00, error = 2.98e-12

real      0m3.299s
user      0m3.298s
sys       0m0.001s

$ time ./comptraprecomp 200000 10
approximation = 1.7182818284620265e+00
exp(1) - 1 = 1.7182818284590451e+00, error = 2.98e-12

real      0m0.743s
user      0m4.003s
sys       0m0.004s
$
```

Tasking with OpenMP

1 Programming Parallel Shared Memory Computers

- tasking with OpenMP

2 Parallel Recursive Functions

- the Fibonacci numbers
- parallel recursive quadrature

3 Bernstein's Conditions

- dependency analysis

4 Task Dependences

- flow dependence
- parallel blocked matrix multiplication

dependency analysis

Which statements can be executed in parallel?

Let u be an operation. Denote:

- $\mathcal{R}(u)$ is the set of memory cells u reads,
- $\mathcal{M}(u)$ is the set of memory cells u modifies.

Two operations u and v are independent if

$$\mathcal{M}(u) \cap \mathcal{M}(v) = \emptyset, \quad \text{and}$$

$$\mathcal{M}(u) \cap \mathcal{R}(v) = \emptyset, \quad \text{and}$$

$$\mathcal{R}(u) \cap \mathcal{M}(v) = \emptyset.$$

The above conditions are known as *Bernstein's conditions*.

Checking Bernstein's conditions is easy for operations on scalars,
is more difficult for array accesses,
and is almost impossible for pointer dereferencing.

an example

Let x be some scalar and consider two statements:

- $u = [x = x + 1],$
- $v = [x = x + 2].$

We see that u and v are independent of each other, because u followed by v or v followed by u is equivalent to $w = [x = x + 3].$

However, execution of u and v happens by a sequence of more elementary instructions:

- $u: r1 = x; r1 += 1; x = r1;$
- $v: r2 = x; r2 += 2; x = r2;$

where $r1$ and $r2$ are registers.

The elementary instructions are no longer independent.

some references

- A. J. Bernstein:
Analysis of Programs for Parallel Processing.
IEEE Transactions on Electronic Computers 15(5):757–763, 1966.
- P. Feautrier: **Bernstein's Conditions.**
In *Encyclopedia of Parallel Computing*, edited by David Padua,
pages 130–133, Springer 2011.
- B. Wilkinson and M. Allen: **Parallel Programming. Techniques and Applications Using Networked Workstations and Parallel Computers.** 2nd Edition. Prentice-Hall 2005.

Tasking with OpenMP

1 Programming Parallel Shared Memory Computers

- tasking with OpenMP

2 Parallel Recursive Functions

- the Fibonacci numbers
- parallel recursive quadrature

3 Bernstein's Conditions

- dependency analysis

4 Task Dependences

- flow dependence
- parallel blocked matrix multiplication

the `depend` clause

The order of execution of tasks can be ordered.

In the `depend` clause, we consider two dependence types:

- 1 The `in` type: the task depends on the sibling task(s) that generates the item followed by the `in:`.
- 2 The `out` type: if an item appeared following an `in:` then there should be task with the clause `out`.

copied from the OpenMP API Examples

```
#include <stdio.h>
#include <omp.h>

int main ( int argc, char *argv[] )
{
    int x = 1;

    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp task shared(x) depend(out: x)
        x = 2;
        #pragma omp task shared(x) depend(in: x)
        printf("x = %d\n", x);
    }
    return 0;
}
```

about the example

In the parallel region, the `single` construct indicates that every instruction needs to be execute only once.

- One task assigns 2 to `x`.
- Another task prints the value of `x`.

Without `depend`, tasks could execute in any order, and the program would have a *race condition*.

The `depend` clauses force the ordering of the tasks.
The example always prints `x = 2`.

Tasking with OpenMP

1 Programming Parallel Shared Memory Computers

- tasking with OpenMP

2 Parallel Recursive Functions

- the Fibonacci numbers
- parallel recursive quadrature

3 Bernstein's Conditions

- dependency analysis

4 Task Dependences

- flow dependence
- parallel blocked matrix multiplication

blocked matrix multiplication

Our last example also comes from the OpenMP API Examples.

Consider the product of two blocked matrices A with B :

$$\begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix}.$$

where

$$C_{i,j} = A_{i,1}B_{1,j} + A_{i,2}B_{2,j},$$

for all i and j .

The arguments of the depend clauses are blocked matrices.

matrices are pointers to rows

Allocating a matrix of dimension `dim`:

```
{  
    double **A;  
    int i;  
  
    A = (double**)calloc(dim, sizeof(double*));  
  
    for(i=0; i<dim; i++)  
        A[i] = (double*)calloc(dim, sizeof(double));  
}
```

Every row `A[i]` is allocated in the loop.

multiplying blocked matrices of random doubles

At the command line, we specify

- 1 the block size, the size of each block,
- 2 the number of blocks in every matrix, and
- 3 the number of threads.

The dimension equals the block size times the number of blocks.

The parallel region:

```
#pragma omp parallel
#pragma omp single
matmatmul(dim,blocksize,A,B,C);
```

One single thread calls the function `matmatmul`.

The `matmatmul` generates a large number of tasks.

the function `matmatmul`

```
void matmatmul
( int N, int BS,
  double **A, double **B, double **C )
{
    int i, j, k, ii, jj, kk;

    for(i=0; i<N; i+=BS)
    {
        for(j=0; j<N; j+=BS)
        {
            for(k=0; k<N; k+=BS)
            {
```

The triple loop computes the block $C_{i,j}$

multiplying each block in one task

Each task has its own indices ii , jj , and kk .

```
#pragma omp task private(ii, jj, kk) \  
    depend(in: A[i:BS][k:BS], B[k:BS][j:BS]) \  
    depend(inout: C[i:BS][j:BS])  
{  
    for(ii=i; ii<i+BS; ii++)  
        for(jj=j; jj<j+BS; jj++)  
            for(kk=k; kk<k+BS; kk++)  
                C[ii][jj] = C[ii][jj] + A[ii][kk]*B[kk][jj];  
}
```

The `inout` dependence type `C[i:BS][j:BS]` expresses that the dependencies of the update of the block $C_{i,j}$

runs with 2 and 4 threads

```
$ gcc -fopenmp -O3 -o matmulomp matmulomp.c
```

```
$ time ./matmulomp 500 2 2
```

```
real    0m0.828s
user    0m1.558s
sys     0m0.020s
```

```
$ time ./matmulomp 500 2 4
```

```
real    0m0.445s
user    0m1.575s
sys     0m0.017s
$
```

parallel linear algebra

PLASMA (Parallel Linear Algebra Software for Multicore Architectures) is a numerical library intended as a successor to LAPACK for solving problems in dense linear algebra on multicore processors.

A. YarKhan, J. Kurzak, P. Luszczek, J. Dongarra:

Porting the PLASMA Numerical Library to the OpenMP Standard.
International Journal of Parallel Programming, May 2016.

exercises

- 1 Label the six elementary operations in the example on slide 16 as $u_1, u_2, u_3, v_1, v_2, v_3$, write for each the sets $\mathcal{R}(\cdot)$ and $\mathcal{M}(\cdot)$.
Based on the dependency analysis, arrange the six instructions for a correct parallel computation.
- 2 The block size, number of blocks, and number of threads are the three parameters in `matmulomp`.
Explore experimentally with `matmulomp` the relationship between the number of blocks and the number of threads.
For which values do you obtain a good speedup?